

A1: Python and sounds

Course on Audio Signal Processing for Music Applications

Introduction

Welcome to the course! This first programming assignment is for you to get familiar with the assignment submission system and to test out some very basic audio operations using Python. There are three exercises in this assignment. 1) Reading an audio file, 2) Basic operations with audio, and 3) Python array indexing.

Before attempting the assignment, please go through the general guidelines for all the programming assignments. You can find them in ‘Programming assignment-Guidelines.pdf’ in the assignment folder (A1) or on the instructions page of the programming assignment 1 (https://class.coursera.org/audio-002/assignment/view?assignment_id=27).

Relevant concepts

Python: Python is a powerful and easy to learn programming language, which is used in a wide variety of application areas. We will use python in all our programming assignments and in this first assignment you will start learning about it by performing some basic operations with sound files. For a quick introduction to python watch the first programming lecture of the first week of the course. (https://class.coursera.org/audio-002/lecture/view?lecture_id=21).

Wav file: The wav file format is a lossless format to store sounds on a hard drive. Each audio sample is stored as a 16 bit (sometimes also as 32 bits) integer number. In this course we will work with only one type of audio files. All the sound files we use in the assignments should be wav files that are mono (one channel), in which the samples are stored in 16 bits, and that use the sampling rate of 44100 Hz. Once read into python, the samples will be converted to floating point values with a range from -1 to 1, resulting in a one-dimensional array of floating point values.

A1-Part-1: Reading a wav audio file (*3 points*)

Complete the function `readAudio(inputFile)` in the file `A1Part1.py` so that it reads an audio file and returns 10 consecutive samples of the file starting from the 50001th sample. This means that the output should exactly contain the 50001th sample to the 50010th sample (10 samples).

The input to the function is the file name (including the path) and the output should be a numpy array containing 10 samples.

If you use the `wavread` function from the `utilFunctions` module the input samples will be automatically converted to floating point numbers with a range from -1 to 1, which is what we want.

Remember that in python, the index of the first sample of an array is 0 and not 1.

If you run your code using `piano.wav` as the input, the function should return the following numpy array with 10 samples: `array([-0.06213569, -0.04541154, -0.02734458, -0.0093997, 0.00769066, 0.02319407, 0.03503525, 0.04309214, 0.04626606, 0.0441908], dtype=float32)`.

```
def readAudio(inputFile):
    """
    Input:
        inputFile: the path to the wav file
    Output:
        The function should return a numpy array that
        contains 10 samples of the audio.
    """
    ## Your code here
```

A1-Part-2: Basic operations with audio (*3 points*)

Complete the function `minMaxAudio(inputFile)` in the file `A1Part2.py` so that it reads an audio file and returns the minimum and the maximum values of the audio samples in that file.

The input to the function is the wav file name (including the path) and the output should be two floating point values returned as a tuple.

If you run your code using `oboe-A4.wav` as the input, the function should return the following output: `(-0.83486432, 0.56501967)`

```
def minMaxAudio(inputFile):
    """
    Input:
        inputFile: file path to the wav file
    Output:
        A tuple of the minimum and the maximum value of the audio
        samples, like: (min_val, max_val)
    """
    ## Your code here
```

A1-Part-3: Python array indexing (*4 points*)

Complete the function `hopSamples(x,N)` in the file `A1Part3.py` so that given a numpy array `x`, the function returns every Nth element in `x`, starting from the first element.

The input arguments to this function are a numpy array `x` and a positive integer `N` such that `N` is less than the number of elements in `x`. The output of this function should be a numpy array.

If you run your code with `x = np.arange(10)` and `N = 2`, the function should return the following output: `array([0, 2, 4, 6, 8])`.

```
def hopSamples(x,N):  
    """  
    Inputs:  
        x: input numpy array  
        N: a positive integer, (indicating hop size)  
    Output:  
        A numpy array containing every Nth element in x, starting  
        from the first element in x.  
    """  
    ## Your code here
```

Grading

The maximum number of points for this assignment is 10.